

QUESTION 1: TensorFlow vs PyTorch Differences

1. Computational Graph Definition:

TensorFlow: Uses static computational graphs. You define the graph first, then execute it.

PyTorch: Uses dynamic computational graphs (define by run)

2. Syntax:

TensorFlow: More declarative style.

PyTorch: More imperative and intuitive, feels like regular Python code.

3. Debugging:

TensorFlow: Debugging static graphs can be challenging.

PyTorch: Easier debugging using standard Python tools like pdb.

4. Deployment:

TensorFlow: Strong production deployment tools (TensorFlow Serving, TensorFlow Lite, TFX)

PyTorch: Historically weaker in production (improving with TorchServe, TorchScript)

when to use tensor flow and pytorch;

TensorFlow: Best for production-ready systems, mobile/edge deployment, or when using Google Cloud tools.

PyTorch: Best for research, prototyping, and projects where flexibility and quick iteration matter most.

QUESTION 2: Jupyter Notebooks Use Cases in AI Development

Use Case 1: Exploratory Data Analysis and Model Prototyping

Data Exploration: Quickly visualize datasets, check distributions, identify outliers

Iterative Model Development: Test different architectures, hyperparameters, and see results immediately

Example: Loading a dataset, creating visualizations, trying different preprocessing techniques, and training multiple small models to gauge performance - all in an interactive environment.

Use Case 2: Educational Demonstrations and Documentation

Teaching Concepts: Combine explanatory text, mathematical equations, and executable code to explain AI concepts

Reproducible Research: Share complete analysis pipelines with results embedded

Example: Creating a tutorial that explains attention mechanisms in transformers with runnable code examples and visualizations of attention weights.

QUESTION3: How spaCy Enhances NLP vs Basic Python String Operations

Basic Python string methods (`split()`, `replace()`, `find()`) treat text as raw characters with no linguistic awareness. spaCy, on the other hand, provides a linguistically informed NLP pipeline

Tokenization & Sentence Segmentation: Splits text into words/sentences with awareness of punctuation, abbreviations, etc.

Part-of-Speech Tagging & Dependency Parsing: Identifies grammatical roles and syntactic relationships.

Named Entity Recognition (NER): Extracts entities like people, places, organizations, dates.

Lemmatization: Reduces words to their base form (e.g., running → run).

Efficiency: Built in Cython, spaCy processes millions of words quickly—far faster than naive Python loops.

QUESTION 3:

Target Applications

Scikit-learn;

Focused on classical machine learning: regression, classification, clustering, dimensionality reduction, and preprocessing.

Great for small to medium-sized datasets and traditional statistical models.

Not designed for deep learning or large-scale neural networks.

TensorFlow;

Built for deep learning and large-scale neural networks.

Supports advanced architectures (CNNs, RNNs, Transformers) and distributed training.

Includes deployment tools for mobile, web, and production environments.

Ease of Use for Beginners

Scikit-learn;

Very beginner-friendly with a *consistent API* (fit(), predict(), transform()).

Minimal boilerplate code—ideal for quickly testing algorithms and pipelines.

Excellent documentation and simple integration with NumPy, Pandas, and Matplotlib.

TensorFlow;

Steeper learning curve, especially for those new to deep learning.

High-level APIs like *Keras* make it easier, but still more complex than Scikit-learn.

Requires understanding of tensors, computational graphs, and GPU acceleration.

Community Support

Scikit-learn;

Mature, stable, and widely used in academia and industry for classical ML.

Strong documentation, tutorials, and a large base of contributors.

Slower pace of change—focused on reliability rather than cutting-edge features.

TensorFlow;

Backed by Google with a massive global community.

Rich ecosystem (TensorFlow Hub, TensorBoard, TensorFlow Lite, TensorFlow.js).

Rapidly evolving, with strong support for state-of-the-art deep learning research and production deployment.

CLASSICAL ML WITH SCIKIT-LEARN

```
First 5 rows of the dataset:
  sepal_length  sepal_width  petal_length  petal_width  species
0           5.1           3.5           1.4           0.2  setosa
1           4.9           3.0           1.4           0.2  setosa
2           4.7           3.2           1.3           0.2  setosa
3           4.6           3.1           1.5           0.2  setosa
4           5.0           3.6           1.4           0.2  setosa

Missing values in each column:
sepal_length    0
sepal_width     0
petal_length    0
petal_width     0
species         0
dtype: int64

Species encoding:
setosa -> 0
versicolor -> 1
virginica -> 2

Features shape: (150, 4)
Target shape: (150,)

Training set shape: (120, 4) (120,)
Testing set shape: (30, 4) (30,)

Model trained successfully.
Predictions made on test set.

Accuracy: 0.93
Precision (macro): 0.93
Recall (macro): 0.93

Classification Report:
      precision    recall  f1-score   support

   setosa         1.00      1.00      1.00        10
```

DEEPLARNING WITH TENSORFLOW

```
1875/1875 — 69s 32ms/step - accuracy: 0.9100 - loss: 0.2950 - val_accuracy: 0.9847 - val_loss: 0.0465
Epoch 2/10
1875/1875 — 76s 32ms/step - accuracy: 0.9863 - loss: 0.0438 - val_accuracy: 0.9882 - val_loss: 0.0353
Epoch 3/10
1875/1875 — 58s 31ms/step - accuracy: 0.9907 - loss: 0.0279 - val_accuracy: 0.9906 - val_loss: 0.0330
Epoch 4/10
1875/1875 — 84s 32ms/step - accuracy: 0.9940 - loss: 0.0186 - val_accuracy: 0.9877 - val_loss: 0.0354
Epoch 5/10
1875/1875 — 57s 31ms/step - accuracy: 0.9952 - loss: 0.0151 - val_accuracy: 0.9875 - val_loss: 0.0407
Epoch 6/10
1875/1875 — 59s 31ms/step - accuracy: 0.9965 - loss: 0.0101 - val_accuracy: 0.9900 - val_loss: 0.0351
Epoch 7/10
1875/1875 — 57s 31ms/step - accuracy: 0.9968 - loss: 0.0081 - val_accuracy: 0.9904 - val_loss: 0.0374
Epoch 8/10
1875/1875 — 58s 31ms/step - accuracy: 0.9975 - loss: 0.0084 - val_accuracy: 0.9907 - val_loss: 0.0356
Epoch 9/10
1875/1875 — 58s 31ms/step - accuracy: 0.9981 - loss: 0.0057 - val_accuracy: 0.9907 - val_loss: 0.0393
Epoch 10/10
1875/1875 — 61s 33ms/step - accuracy: 0.9978 - loss: 0.0058 - val_accuracy: 0.9921 - val_loss: 0.0381
313/313 — 3s 10ms/step - accuracy: 0.9920 - loss: 0.0386
```

Test accuracy: 0.9921

```
1/1 — 0s 122ms/step
1/1 — 0s 58ms/step
1/1 — 0s 49ms/step
1/1 — 0s 58ms/step
1/1 — 0s 54ms/step
```

True Label: 2, Predicted Label: 2



True Label: 8, Predicted Label: 8



True Label: 2, Predicted Label: 2



True Label: 5, Predicted Label: 5



True Label: 2, Predicted Label: 2



NLP WITH SPACY

```
print("\n--- Extracted Entities ---")
for ent in doc.ents:
    if ent.label_ in ['PRODUCT', 'ORG']:
        print(f"Entity: {ent.text}, Type: {ent.label_}")

print("\n--- Sentiment Analysis Results ---")
print(f"Overall sentiment: {sentiment}")
print(f"Sentiment score: {sentiment_score}")
```

```
--- Extracted Entities ---
Entity: Samsung, Type: ORG
Entity: Apple, Type: ORG
```

```
--- Sentiment Analysis Results ---
Overall sentiment: positive
Sentiment score: 4
```